

January 4, 2026

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>System Architecture</b>	<b>2</b>
<b>3</b>	<b>Usage</b>	<b>3</b>
3.1	Running the Simulation . . . . .	3
3.2	Command-Line Arguments . . . . .	3
3.3	Input and Output . . . . .	4
<b>4</b>	<b>How to Add a New Distribution</b>	<b>5</b>
4.1	Step 1: Define the Input File Format . . . . .	5
4.2	Step 2: Create a Parser Function in <code>distros.py</code> . . . . .	5
4.3	Step 3: Implement the Sigma Function in <code>distros.py</code> . . . . .	5
4.4	Step 4: Register in <code>main.py</code> . . . . .	6

# 1 Introduction

To achieve high performance, the simulation logic is: This document describes a high-performance tool designed for the Multiserver-Job Queueing Model (MJQM) [1, 2]. The tool uses the Sub-Perfect Sample algorithm presented in [3] to find the stationary distribution of the system’s workload.

To achieve high performance, the simulation logic is:

- GPU-Accelerated: core numerical operations are vectorized and executed on the GPU using the `cupy` library.
- Batch-Processed: the simulation runs thousands of independent trials (`Ntries`) simultaneously in large batches to maximize GPU utilization.
- Modular: the architecture is designed to easily swap out or add new service time distributions, which is the primary focus of this documentation.

The core of the simulation is a backward-simulation approach based on a Lindley-like recursion, which is implemented in the `one_sampling` function. The perfect sampling is achieved by a “doubling” scheme (starting at length  $L$  and doubling to  $2L$ ,  $4L$ , ...) until the initial workload state no longer affects the final workload state.

# 2 System Architecture

The tool is built on three main Python files.

`main.py` This is the main entry point. It parses command-line arguments and acts as a controller. It selects the correct service time distribution based on the `-dist` argument. It selects the appropriate parsing function accordingly, and the corresponding service time generator function and passes them to the simulation engine.

`perfsmpl.py` This file is the core of the whole program, and it is the distribution-agnostic simulation engine.

- `perfect_sampling()`: manages the overall simulation, including batching, the doubling-convergence loop, and writing results to the CSV file.
- `generate_chunk()`: generates blocks of random numbers on the CPU (for robust seeding) and then calls the specific `compute_sigmas` function (passed in from `main.py`) to convert these uniform randoms into service times on the GPU.
- `one_sampling()`: performs one vectorized step of the backward simulation. It is completely general and only requires the final `sigmas`, `taus`, and `alphas`.

`distros.py` This file is a library for the different possible distributions. For each supported probability distribution, it must provide two functions:

1. A **parser function** (e.g., `exp_parser`) that knows how to read a line from an input text file and return a dictionary of parameters.
2. A **sigma function** (e.g., `compute_sigmas_exp`) that takes uniform random numbers and generates service times on the GPU according to that distribution.

## 3 Usage

### 3.1 Running the Simulation

The simulation is launched from `main.py`. A `run.sh` script is provided to automate the process of running experiments for different distributions and arrival rates (`lambda`).

```
1 #!/bin/bash
2
3 # Fixed Parameters
4 nsamples_values=(100000)
5 N=20
6 L=100
7 split_size=100
8 batch_size=10000
9 device="cuda:0"
10 seed=42
11
12 # Define Distributions
13 declare -A dist_input
14 declare -A dist_lambdas
15
16 # Exponential
17 dist_input["exp"]="synth_20_exp"
18 dist_lambdas["exp"]="1.3338 1.5244 1.7149"
19
20 # Erlang-k
21 dist_input["erlk"]="synth_20_erlk"
22 dist_lambdas["erlk"]="1.2515 1.4303 1.6091 1.7832"
23
24 # Loop Over Distributions
25 for dist in "${!dist_input[@]}"; do
26     in_fname="${dist_input[$dist]}"
27     IFS=' ' read -r -a lambdas <<< "${dist_lambdas[$dist]}"
28
29     for lambda in "${lambdas[@]}"; do
30         for nsamples in "${nsamples_values[@]}"; do
31             echo ">>> Running with lam = $lambda"
32             time python3 main.py \
33                 -dist "$dist" \
34                 -device "$device" \
35                 -in_fname "$in_fname" \
36                 -N "$N" \
37                 -lam "$lambda" \
38                 -nsamples "$nsamples" \
39                 -L "$L" \
40                 -split_size "$split_size" \
41                 -batch_size "$batch_size" \
42                 -seed "$seed"
43         done
44     done
45 done
```

Listing 1: Example `run.sh` script

### 3.2 Command-Line Arguments

The `main.py` script accepts the following arguments.

Table 1: Command-Line Arguments

Argument	Type	Description
<code>-dist</code>	str	<b>Required.</b> The distribution to use. Choices: <code>exp</code> , <code>bpar</code> , <code>erlk</code> , <code>hyperexp</code> .
<code>-device</code>	str	The GPU device to use (e.g., <code>cuda:0</code> ).
<code>-in_fname</code>	str	Name of the input parameter file (e.g., <code>synth_20_exp</code> ) located in the <code>inputs/</code> folder.
<code>-N</code>	int	Number of servers in the system.
<code>-lam</code>	float	The overall arrival rate (lambda).
<code>-nsamples</code>	int	Total number of simulation trials to run.
<code>-L</code>	int	The <i>initial</i> simulation length for the doubling algorithm.
<code>-split_size</code>	int	The chunk size ( $L$ ) for <code>one_sampling</code> . A large <code>-L</code> is split into multiple chunks of this size to avoid GPU OOM errors.
<code>-batch_size</code>	int	The number of trials to run in a single batch. <code>-nsamples</code> is split into batches of this size to avoid CPU/GPU OOM errors.
<code>-seed</code>	int	The master seed for the random number generator to provider reproducibility.

### 3.3 Input and Output

- **Inputs:** All parameter input files (e.g., `synth_20_exp.txt`) are expected to be in a directory named `inputs/`.
- **Outputs:** All simulation results are saved as `.csv` files in a directory named `results/`.

## 4 How to Add a New Distribution

This is the main purpose of the tool's design. To add a new service time distribution (e.g., the Log-Normal distribution), you must complete four steps.

### 4.1 Step 1: Define the Input File Format

First, decide what parameters your distribution needs for each class. Create a `.txt` file in the `inputs/` directory.

For our Log-Normal example, let's say we need *mean* and *variance* of the underlying Normal distribution. The format might be: `(size, prob, mu, sigma_sq)`. A file `inputs/synth_20_lognormal.txt` might look like:

```
(5, 0.5, 0.1, 0.2)
(10, 0.5, 0.5, 0.3)
```

### 4.2 Step 2: Create a Parser Function in `distros.py`

Add a new parser function to `distros.py` that can read one line of your new file. It must return a dictionary.

```
1 def lognormal_parser(values):
2     """
3         Parses a line for Log-Normal distribution input parameters.
4         Expected format: (size, prob, mu, sigma_sq)
5     """
6     return {
7         'sizes': int(values[0]),
8         'probs': float(values[1]),
9         'mus': float(values[2]),
10        'sigma_sqs': float(values[3])
11    }
```

Listing 2: New parser function in `distros.py`

### 4.3 Step 3: Implement the Sigma Function in `distros.py`

This is the most important step. You must write a GPU-accelerated function that generates the service times.

**Function Signature:** The function **must** have the signature:

```
def compute_sigmas_lognormal(rng_sigma, class_indices, **params):
```

- `rng_sigma`: This is a **list of CuPy arrays** containing uniform random numbers between  $[0, 1]$ . The number of arrays in the list is determined by `num_streams` set in `main.py`.
- `class_indices`: This is a 2D CuPy array ( $B, L$ ) that maps each arrival to its class ID.
- `**params`: This is a dictionary containing the parameter arrays loaded by your parser (e.g., `params['mus']`, `params['sigma_sq']`).

**Example Implementation:** To generate a Log-Normal variate, we can use two uniform randoms ( $U_1, U_2$ ) with the Box-Muller transform to get a standard normal variate  $Z$ , and then transform it. This means we will need **2 random number streams**.

```

1 import cupy as cp
2
3 def compute_sigmas_lognormal(rng_sigma, class_indices, **params):
4     """
5         Generates Log-Normal distributed service times (sigmas) on the GPU.
6         Uses the Box-Muller transform, requiring 2 random streams.
7     """
8
9     # --- 1. Get Random Streams ---
10    # We need two independent uniform random streams
11    unif_1 = rng_sigma[0]
12    unif_2 = rng_sigma[1]
13
14    # --- 2. Get Class-Specific Parameters ---
15    # Use class_indices to "map" arrivals to their parameters
16    mus = params['mus'][class_indices]
17    sigma_sqs = params['sigma_sq'][class_indices]
18    sigmas_norm = cp.sqrt(sigma_sqs) # StDev of the normal
19
20    # --- 3. Implement the Distribution (Inverse Transform, etc.) ---
21    # Box-Muller transform to get a standard normal variate (Z)
22    # We only need one of the two generated variates.
23    Z = cp.sqrt(-2.0 * cp.log(unif_1)) * cp.cos(2.0 * cp.pi * unif_2)
24
25    # --- 4. Scale and Shift ---
26    # Transform the standard normal (Z) to N(mu, sigma_sq)
27    log_Y = mus + Z * sigmas_norm
28
29    # Exponentiate to get the Log-Normal variate
30    sigmas_final = cp.exp(log_Y)
31
32    return sigmas_final

```

Listing 3: New sigma function in distros.py

#### 4.4 Step 4: Register in main.py

Finally, add an `elif` block to `main.py` to tell the controller about your new "lognormal" strategy.

```

1 # ... other elif blocks ...
2 elif args.dist == 'erlk':
3     parser = distros.erlk_parser
4     in_data, nClasses = distros.load_params(args.in_fname, parser)
5     num_streams = int(cp.max(in_data['ks']))
6     compute_sigmas_foo = distros.compute_sigmas_erlk
7     distro_label = 'erlk'
8
9 # --- ADD YOUR NEW BLOCK HERE ---
10 elif args.dist == 'lognormal':
11     parser = distros.lognormal_parser # <-- Your new parser
12     in_data, nClasses = distros.load_params(args.in_fname, parser)
13     num_streams = 2                  # <-- We need 2 for Box-Muller
14     compute_sigmas_foo = distros.compute_sigmas_lognormal # <-- Your
15     new func
16     distro_label = 'lognormal'
17
18 # --- END OF NEW BLOCK ---

```

```
19     print(f"Error: Unknown distribution '{args.dist}'")
```

Listing 4: Modification to `main.py`

After these four steps, you can run the simulation by passing the new argument:

```
python3 main.py -dist "lognormal" -in_fname "synth_20_lognormal" ...
```

## References

- [1] M. Harchol-Balter, “The Multiserver Job Queueing Model,” *Queueing Systems: Theory and Applications*, vol. 100, pp. 201–203, 2022.
- [2] D. Olliari, M. Ajmone Marsan, S. Balsamo, and A. Marin, “The saturated Multiserver Job Queueing Model with two classes of jobs: Exact and approximate results,” *Performance Evaluation*, vol. 162, p. 102370, 2023.
- [3] F. Baccelli, D. Olliari, M. Ajmone Marsan, and A. Marin, “The Multiserver-Job Stochastic Recurrence Equation for Cloud Computing Performance Evaluation,” *pomacs wannabe*, hopefully 2026.